

There are two problems this week, but both are dealing with the same set of data. In the repository is a script to recreate a database of airplane flights over the month of January in 2020. If you are interested in more, similar type data, I grabbed the original data from the Bureau of Transportation Statistics [here](#).

As per usual, you should follow the link here to accept the assignment and get access to the repository:

Assignment link: <https://classroom.github.com/a/4-DI8wYv>

In the repository lives a zip file named `hw6.zip`. You should download that zip file and extract it, which has inside two SQL files (compressing made them small enough to upload to GitHub). These `.sql` files are a bit different from what you have seen in the past though, as they are technically a *database dump*. These are a common way to generate exact duplicates of the contents of a database. *However*, interacting with them can have some differences. The easiest, and fastest, method is to use the command line (terminal) tool `psql` that should have been installed along with the rest of PostgreSQL at the start of the semester. If you open a terminal (or command prompt) at the location of the `hw6_fast.sql` file (or open a terminal anywhere and then navigate to that location), the command that you want to run is simply:

```
psql -d analysis -U yourusername -f hw6_fast.sql
```

where `analysis` is presumably the name of the database where you've been placing everything, and `yourusername` is whatever user name you've been using with your database. If you happened to set up the database server on a non-default port, you'll also need to include `-p portnum`. If all goes successfully, this will create a new `hw6` schema in your `analysis` database and set up the 6 different tables, complete with data, constraints, and primary and foreign keys. You should be able to confirm this by opening whatever client you usually use and querying the tables.

If this is not working, for whatever reason, I've tried to give you another option as well. The `hw6_slow.sql` file is also a database dump, but one which has been tweaked to use standard `INSERT INTO` statements. The result is that you should be able to run this file like any other collection of SQL commands in your client, but the price is that it will be MUCH slower to populate the initial tables. I'd recommend the first option if possible, but if you are having absolutely no luck getting `psql` to work in the terminal, I think this will work. *If you are still having trouble getting the tables generated, please reach out and I'm happy to help you through this!*

The benefit of loading a database dump, of course, is that the database structure is all ready for you, with a central table `flights` having information about individual flights as well as references to other tables with information about airports, carriers, and city markets (to name a few). An image of the ERD is included on the next page, as well as in the repository, for your reference, or you can look at the version online [here](#).

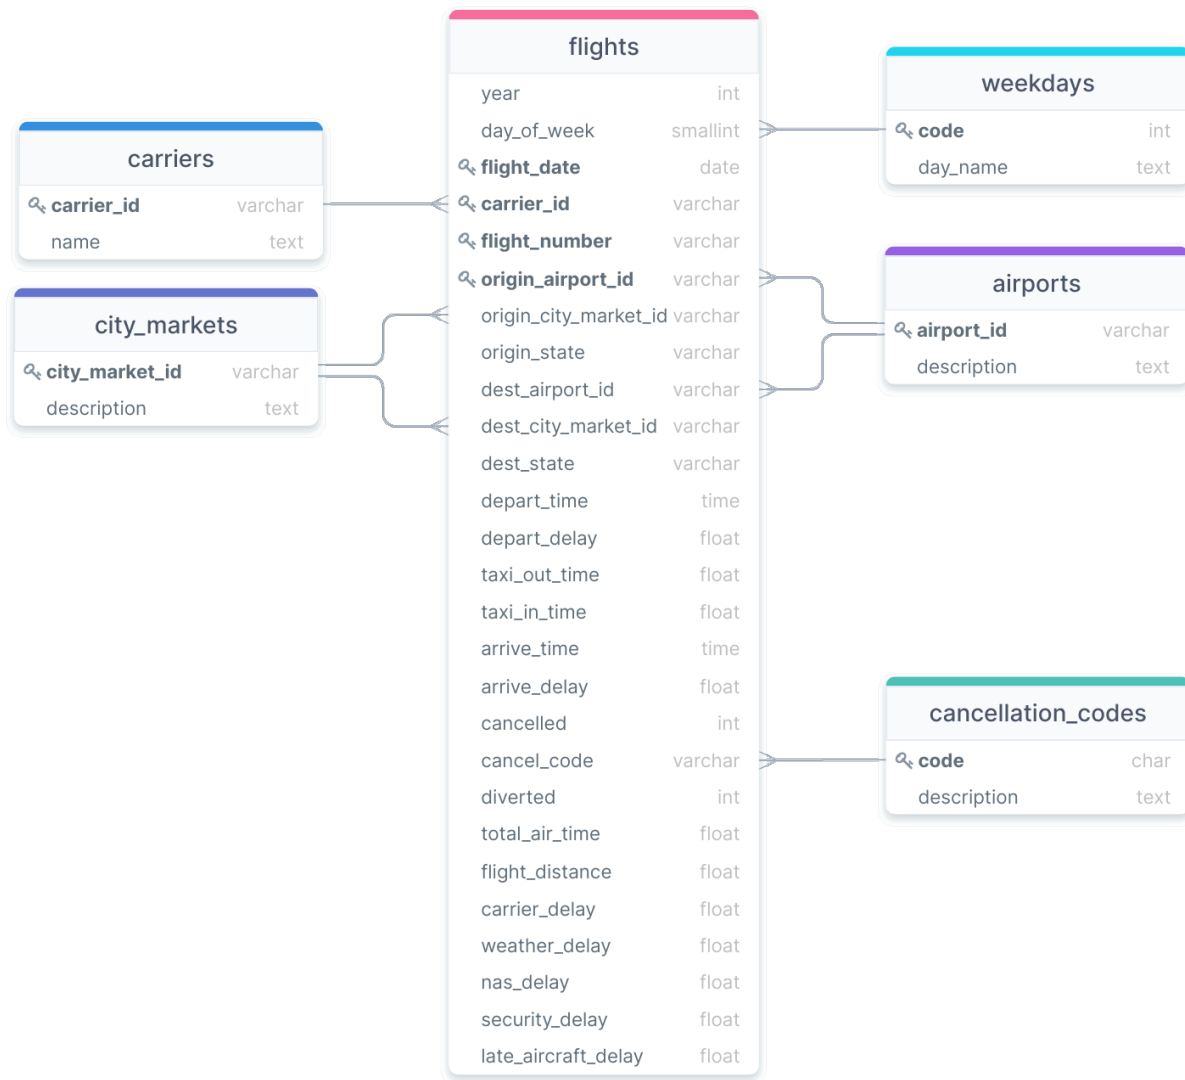


Figure 1: Entity Relationship Diagram of the tables and relationships included in the hw6 database. The **flights** table is the main table, and references several smaller tables for extra details (mainly names) of carriers, airports, and city markets (the metropolitan area that a particular airport serves).

1. This initial batch of questions will focus on being able to use keywords such as **GROUP BY** effectively. For each you should **include any queries you used** to determine the answer, as well as the answer itself. *You may need several queries or temporary tables to answer some of the questions. Show all your work.*
  - (a) (4 points) What is the name of the airport which, over the given month's data, has the greatest average delay in arriving flights? Limit your answer to only include airports with at least 100 arriving flights over the course of the month.
  - (b) (4 points) Which 3 US cities have the greatest number of inbound flights across all airports serving that city? (Some cities have multiple airports, which is what the `city_market_id` is supposed to assist with.)
  - (c) (4 points) Which 5 airports have the greatest number of outgoing flights *per day* on average? How many flights (on average) are leaving those airports each day?
  - (d) (6 points) For each major airline carrier, which two airports represent the endpoints of that carrier's longest flight (by distance)? There are 17 distinct carriers in the month represented in the dataset, so your output table should also have 17 rows with the name of the carrier and the name of both airports. Dealing with round trip duplicates here can be a little tricky, so think about how you could reliably filter out one of the directions. It doesn't matter whether you choose PDX to SEA or SEA to PDX (for instance), you just need one of the pairings for each carrier.
2. Just looking at the numeric outputs from our queries isn't always the most enlightening. Sometimes a little data visualization goes a long way. Depending on what you have done in the past, you may have a host of data visualization tools in your toolbox, ranging from R through Python to Excel. Any are completely appropriate for the following questions, but regardless of what software you use, ensure that your charts have clearly labeled axes and a title describing what they depict. You may need to save the output of your query in a format that your visualization software can understand (frequently CSV). Remember that you already *know* how to do this, and it is very similar to how you import information into a table.
  - (a) (4 points) Create a simple bar chart showing the total number of flights that left airports each day of the week.
  - (b) (6 points) This second part is more open-ended. Define a question that you are interested in investigating in this current data set. There is lots of potential data here that I haven't already asked questions about that might be interesting! The only requirement is that you must use **GROUP BY** in the course of answering your question, and you must include a visualization to help you answer the question. Clearly state both your question, your work to answer the question, and your visualization in your submission.